

NLP ASSIGNMENT 2 REPORT

NAME: SRUJAN ALUMULA

UID: U42726372

▪ **Introduction: Objectives and overview of the work.**

Analysing and contrasting the effectiveness of character-level and word-level tokenization techniques on a text classification task using the IMDB dataset is the main goal of this work.

Experimenting with different hyperparameters, such as activation functions, learning rates, batch sizes, and optimizers, is also intended to maximize the performance of the model.

Tokenization plays a crucial role in the text preprocessing significantly impacting the model performance. This study compares the two tokenization approaches to understand their influence on the model's ability to classify the text data.

Hyperparameter optimization is used systematically to determine the best combination of parameters for achieving the better high accuracy.

▪ **Methodology: Detailed explanation of tokenization changes and hyper-parameter optimization strategy.**

Tokenization Changes :Character Level vs Word Level

The main focus was on two distinct types that are character level tokenization and word level tokenization. The main objective was to compare their impact on the model performance, training efficiency and the final evaluation metrics.

Character Level Tokenization

Character level tokenization involves breaking down text into the individual characters, rather than the whole words. This approach is especially useful for handling rare words, misspellings, or language with complex word structures. In this project the following steps were followed for character level tokenization.

Text Preprocessing : The text data first preprocessed to convert all the characters to lowercase ensuring the consistency across inputs.

Tokenizer Setup: A character level tokenizer was built using the tensorflow tokenizer api with the charlevel = true parameter. This tokenizer was then fitted on the training data to learn the vocabulary of characters.

Bag of Characters – After Tokenization the text was transformed into the binary bag of characters representation using the matrix. This matrix represented the presence or absence of each character in the vocabulary.

Hyperparameters Optimized:

Activation Function : activation function is the how the weighted sum of inputs is transformed into an output. Several activation functions were tested

Leaky relu – A variant of relu that allows small nonzero gradients when the input is negative potentially preventing dead neurons.

Relu : A widely used activation function that outputs zero for negative inputs and the same value for the positive inputs.

Tanh : A squashed activation function with outputs between -1 and 1 often used the data distribution is centered around 0.

Optimizer : The optimizer is responsible for the adjusting the weights during training . The following optimizers were I tested.

Adam: An adaptive optimizer that computes adaptive learning rated for each parameter.

SGD: Stochastic Gradient Descent, a simple optimizer often used with momentum for faster convergence.

RMSprop: An optimizer that adjusts the learning rate for each parameter based on the recent gradient information.

- And Also the remaining all I tested with these
Learning Rate: [0.001, 0.0005, 0.0001]
- **Hidden Layers:** [1, 2, 3]
- **Hidden Sizes:** [128, 256, 512]
- **Batch Sizes:** [32, 64, 128]
- **Optimizers:** [Adam, SGD, RMSProp]
- **Activation Functions:** [ReLU, Tanh, LeakyReLU]

SEED	Activation Function	Optimizer	Batch Size	Learning Rate	Testing Accuracy	Test Loss
655	Relu	RMSProp	512	0.0001	60.37	66.35
655	Relu	RMSProp	128	0.0001	60.78	66.06
655	Relu	SGD	128	0.0001	53.39	69.07
655	Tanh	Adam	512	0.0001	60.44	66.12
655	LeakyRelu	Adam	512	0.0001	60.44	66.25
655	Relu	Adam	512	0.0001	60.64	66.28
655	Relu	Adam	256	0.001	60.97	65.77
655	Relu	Adam	256	0.0005	60.65	66.11

655	Relu	Adam	256	0.0001	60.75	66.06
655	Relu	Adam	256	0.001	60.66	66.14
655	Relu	Adam	256	0.001	60.64	66.08
655	Relu	Adam	256	0.0001	60.51	66.15
655	Relu	Adam	256	0.0005	60.88	65.90
655	Tanh	Adam	128	0.0001	60.44	66.12
655	Relu	Adam	128	0.001	60.64	66.14
655	Relu	SGD	512	0.0001	60.22	70.54
655	Tanh	Adam	512	0.001	60.68	65.95
655	LeakyRelu	Adam	512	0.001	61.01	65.79
655	LeakyRelu	SGD	512	0.01	53.90	68.91
655	LeakyRelu	RMSProp	512	0.01	56.90	68.41
655	LeakyRelu	Adam	256	0.01	61.04	65.94
655	LeakyRelu	Adam	128	0.01	60.64	66.11
1567	LeakyRelu	Adam	256	0.01	58.74	60.81
45	LeakyRelu	Adam	256	0.01	61.14	65.83

655 (RANDOM) model	LeakyRelu	Adam	256	0.01	58.52	67.29
45	LeakyRelu	Adam	256	0.01	59.10	67.07

Observation :

Among the activation function tested the LeakyRelu is performed consistently well achieving the highest accuracy of 61.16% with a test loss of 65.83 at 256 batch size and learning rate of 0.01.

Optimizer : Adam appeared to be the most reliable optimizer across various configurations producing higher accuracy results compared to RMSProp or SGD.

Batch Size : Larger batches generally resulted in lower accuracy and higher loss compared to medium batches sizes with the leakyrelu and adam optimizer.

Learning Rate: Learning rate around 0.01 consistently showed better results than smaller rates with the higher learning rate tending to achieve slightly better performance.

This Is the Best I verified from the above all

with seed 655 , activation function Leaky Relu , Optimizer Adam , batch Size =256 Learning Rate 0.01

Loading IMDB dataset...

Train samples: 20000, Validation samples: 5000, Test samples: 25000

Tokenizer vocabulary size: 134

Starting training...

Epoch 01	Training Loss: 0.6738	Val Loss: 0.6652	Accuracy: 0.6022	Precision: 0.5959	Recall: 0.5573
Epoch 02	Training Loss: 0.6623	Val Loss: 0.6667	Accuracy: 0.5962	Precision: 0.5620	Recall: 0.7574
Epoch 03	Training Loss: 0.6602	Val Loss: 0.6612	Accuracy: 0.6070	Precision: 0.5845	Recall: 0.6547
Epoch 04	Training Loss: 0.6591	Val Loss: 0.6618	Accuracy: 0.6074	Precision: 0.5771	Recall: 0.7116
Epoch 05	Training Loss: 0.6560	Val Loss: 0.6609	Accuracy: 0.6096	Precision: 0.6067	Recall: 0.5536
Epoch 06	Training Loss: 0.6535	Val Loss: 0.6627	Accuracy: 0.6096	Precision: 0.5818	Recall: 0.6922
Epoch 07	Training Loss: 0.6531	Val Loss: 0.6598	Accuracy: 0.6090	Precision: 0.5959	Recall: 0.6011
Epoch 08	Training Loss: 0.6505	Val Loss: 0.6577	Accuracy: 0.6094	Precision: 0.5905	Recall: 0.6337
Epoch 09	Training Loss: 0.6486	Val Loss: 0.6581	Accuracy: 0.6128	Precision: 0.5985	Recall: 0.6118
Epoch 10	Training Loss: 0.6468	Val Loss: 0.6636	Accuracy: 0.6038	Precision: 0.5710	Recall: 0.7347

Evaluating on test set...

Test Loss: 0.6594 | Test Accuracy: 0.6104 | Test Precision: 0.5877 | Test Recall: 0.7402

These Model is verified with Seed Numbers

Test 2 with - Seed Number 1567

Loading IMDB dataset...
Train samples: 20000, Validation samples: 5000, Test samples: 25000
Tokenizer vocabulary size: 134

Starting training...

Epoch 01		Training Loss: 0.6775		Val Loss: 0.6645		Accuracy: 0.6064		Precision: 0.5776		Recall: 0.7001
Epoch 02		Training Loss: 0.6633		Val Loss: 0.6629		Accuracy: 0.6066		Precision: 0.5803		Recall: 0.6811
Epoch 03		Training Loss: 0.6603		Val Loss: 0.6608		Accuracy: 0.6078		Precision: 0.5848		Recall: 0.6584
Epoch 04		Training Loss: 0.6584		Val Loss: 0.6625		Accuracy: 0.6066		Precision: 0.5761		Recall: 0.7133
Epoch 05		Training Loss: 0.6570		Val Loss: 0.6609		Accuracy: 0.6094		Precision: 0.5819		Recall: 0.6902
Epoch 06		Training Loss: 0.6580		Val Loss: 0.6717		Accuracy: 0.5882		Precision: 0.5520		Recall: 0.7987
Epoch 07		Training Loss: 0.6558		Val Loss: 0.6637		Accuracy: 0.6032		Precision: 0.5706		Recall: 0.7335
Epoch 08		Training Loss: 0.6533		Val Loss: 0.6642		Accuracy: 0.6020		Precision: 0.5687		Recall: 0.7409
Epoch 09		Training Loss: 0.6501		Val Loss: 0.6590		Accuracy: 0.6046		Precision: 0.5976		Recall: 0.5644
Epoch 10		Training Loss: 0.6505		Val Loss: 0.6616		Accuracy: 0.5998		Precision: 0.5696		Recall: 0.7137

Evaluating on test set...

Test Loss: 0.6585 | Test Accuracy: 0.6081 | Test Precision: 0.5874 | Test Recall: 0.7264

Test 3 With seed Number – 45

Loading IMDB dataset...
Train samples: 20000, Validation samples: 5000, Test samples: 25000
Tokenizer vocabulary size: 134

Starting training...

Epoch 01		Training Loss: 0.6750		Val Loss: 0.6663		Accuracy: 0.5994		Precision: 0.5666		Recall: 0.7393
Epoch 02		Training Loss: 0.6616		Val Loss: 0.6638		Accuracy: 0.6016		Precision: 0.5710		Recall: 0.7170
Epoch 03		Training Loss: 0.6603		Val Loss: 0.6602		Accuracy: 0.6106		Precision: 0.5907		Recall: 0.6407
Epoch 04		Training Loss: 0.6579		Val Loss: 0.6602		Accuracy: 0.6112		Precision: 0.5837		Recall: 0.6902
Epoch 05		Training Loss: 0.6572		Val Loss: 0.6600		Accuracy: 0.6060		Precision: 0.5924		Recall: 0.6007
Epoch 06		Training Loss: 0.6587		Val Loss: 0.6595		Accuracy: 0.6088		Precision: 0.5872		Recall: 0.6498
Epoch 07		Training Loss: 0.6537		Val Loss: 0.6593		Accuracy: 0.6068		Precision: 0.5979		Recall: 0.5767
Epoch 08		Training Loss: 0.6523		Val Loss: 0.6578		Accuracy: 0.6046		Precision: 0.5792		Recall: 0.6741
Epoch 09		Training Loss: 0.6508		Val Loss: 0.6597		Accuracy: 0.6064		Precision: 0.5753		Recall: 0.7186
Epoch 10		Training Loss: 0.6476		Val Loss: 0.6580		Accuracy: 0.6078		Precision: 0.5890		Recall: 0.6320

Evaluating on test set...

Test Loss: 0.6583 | Test Accuracy: 0.6114 | Test Precision: 0.6054 | Test Recall: 0.6398

Now These Model is compared with the Random model

Random MLP on IMDB Dataset with seed number of 655

Learning Rate of 0.001 and actiation function leakyrelu with an optimizer ADAM , batach size=256

Loading IMDB dataset...
Train samples: 20000, Validation samples: 5000, Test samples: 25000
Tokenizer vocabulary size: 134

Starting training...

Epoch 01		Training Loss: 0.6927		Val Loss: 0.6886		Accuracy: 0.5420		Precision: 0.5235		Recall: 0.6167
Epoch 02		Training Loss: 0.6844		Val Loss: 0.6839		Accuracy: 0.5552		Precision: 0.5424		Recall: 0.5272
Epoch 03		Training Loss: 0.6811		Val Loss: 0.6827		Accuracy: 0.5668		Precision: 0.5411		Recall: 0.7001
Epoch 04		Training Loss: 0.6786		Val Loss: 0.6792		Accuracy: 0.5776		Precision: 0.5572		Recall: 0.6271
Epoch 05		Training Loss: 0.6765		Val Loss: 0.6784		Accuracy: 0.5766		Precision: 0.5502		Recall: 0.6943
Epoch 06		Training Loss: 0.6751		Val Loss: 0.6767		Accuracy: 0.5822		Precision: 0.5554		Recall: 0.6927
Epoch 07		Training Loss: 0.6736		Val Loss: 0.6750		Accuracy: 0.5860		Precision: 0.5611		Recall: 0.6704
Epoch 08		Training Loss: 0.6726		Val Loss: 0.6738		Accuracy: 0.5880		Precision: 0.5650		Recall: 0.6531
Epoch 09		Training Loss: 0.6717		Val Loss: 0.6731		Accuracy: 0.5886		Precision: 0.5656		Recall: 0.6526
Epoch 10		Training Loss: 0.6711		Val Loss: 0.6732		Accuracy: 0.5874		Precision: 0.5599		Recall: 0.6955

Evaluating on test set...

Test Loss: 0.6729 | Test Accuracy: 0.5882 | Test Precision: 0.5729 | Test Recall: 0.6933

Random MLP on IMDB Dataset with seed number of 45

Learning Rate of 0.001 and activation function leakyrelu with an optimizer ADAM , batch size=256

Loading IMDB dataset...

Train samples: 20000, Validation samples: 5000, Test samples: 25000

Tokenizer vocabulary size: 134

Starting training...

Epoch 01		Training Loss: 0.6918		Val Loss: 0.6869		Accuracy: 0.5468		Precision: 0.5383		Recall: 0.4575
Epoch 02		Training Loss: 0.6852		Val Loss: 0.6817		Accuracy: 0.5714		Precision: 0.5671		Recall: 0.4901
Epoch 03		Training Loss: 0.6806		Val Loss: 0.6793		Accuracy: 0.5844		Precision: 0.5626		Recall: 0.6415
Epoch 04		Training Loss: 0.6772		Val Loss: 0.6790		Accuracy: 0.5746		Precision: 0.5448		Recall: 0.7455
Epoch 05		Training Loss: 0.6750		Val Loss: 0.6757		Accuracy: 0.5874		Precision: 0.5599		Recall: 0.6960
Epoch 06		Training Loss: 0.6732		Val Loss: 0.6745		Accuracy: 0.5904		Precision: 0.5628		Recall: 0.6947
Epoch 07		Training Loss: 0.6718		Val Loss: 0.6719		Accuracy: 0.5900		Precision: 0.5755		Recall: 0.5883
Epoch 08		Training Loss: 0.6707		Val Loss: 0.6714		Accuracy: 0.5924		Precision: 0.5729		Recall: 0.6254
Epoch 09		Training Loss: 0.6699		Val Loss: 0.6722		Accuracy: 0.5918		Precision: 0.5634		Recall: 0.7021
Epoch 10		Training Loss: 0.6696		Val Loss: 0.6722		Accuracy: 0.5914		Precision: 0.5614		Recall: 0.7182

Evaluating on test set...

Test Loss: 0.6707 | Test Accuracy: 0.5910 | Test Precision: 0.5733 | Test Recall: 0.7120

Conclusion:

Through the series of experiments I observed that certain configuration consistently yielded better performance. Specifically, LeakyRelu combined with the Adam optimizer provided the most stable and highest performance across various batch size and learning rate. Among the batch sizes with the learning rate consistently resulted in better accuracy and lower loss. The performance of the model was evaluated using the character level tokenization with the final model performance indicating that word level tokenization provided a more balanced.

The hyperparameter optimization strategy by adjusting activation functions optimizers and learning rate highlighted that Adam was the most effective optimizer, outperforming both RMSProp and SGD. In conclusion the experiments demonstrated the importance of selecting the right combination of hyperparameters as well as the choice between the tokenization in improving the model performance.